

1.

```
# algorithm in python
# assuming the size of array A is greater than 0
def local_min(A: [int]) -> int:
    # base case: if the size of array is 1, return the only element
    # if the size is 2, return the smaller element
    if len(A) == 1:
        return A[0]
    if len(A) == 2:
        return min(A)

    # choose the middle element in the array as the pivot
    pivot = len(A) // 2
    # if the pivot is already a local minimum, return it
    if A[pivot-1] >= A[pivot] <= A[pivot+1]:
        return A[pivot]
    # else, do recursive call on the partition of array with the smaller
    # element, not including the pivot. Since that partition is guaranteed
    # to have a local minimum
    if A[pivot-1] < A[pivot+1]:
        # since end index is not included in list slicing
        return local_min(A[:pivot])
    else:
        return local_min(A[pivot+1:])
```

In the algorithm above, every step other than the recursive call takes constant time; and in the recursive call, the size of input is approximately half of original input size.

Therefore, the recurrence of this algorithm is $T(n) = T(n/2) + 1$, and according to the Master Theorem, $T(n)$ is $\Theta(\log n)$. Additionally, we know that the pivot may already be a local minimum and we may terminate the search early, so the overall asymptotic runtime is $O(\log n)$.

We know the base cases are correct; after checking the base cases, the array will have at least 3 elements so the comparisons with the pivot must be possible. In the partition step, using the partition with the smaller element guarantees that a local minimum will exist in that partition, so we do not need to check the other partition.

2.

algorithm in python

```
def find_median(vector1, vector2, n):
    # base case: if n is small (we chose 2 here), simply use a brute force algorithm
    if n <= 2:
        union = sorted(vector1 + vector2)
        # size of union vector will be either 2 or 4, depending on the value of n
        size = len(union)
        # the median of the two (small) vectors is the average of
        # the middle two elements
        return (union[size/2] + union[size/2-1]) / 2

    # find the median of each vector. Since we know that the two vectors are sorted,
    # the median is simply the middle element
    median1 = len(vector1) // 2
    median2 = len(vector2) // 2
    if n % 2 == 0:
        # if n is even, need to shift the index in one vector by one so that,
        # when we divide the vectors, their size, n, is still equal
        median2 -= 1

    if vector1[median1] < vector2[median2]:
        # in this case, the median will be in the following partitions:
        # vector1: from index of median1 to the end of vector
        # vector2: from beginning to index of median2
        # divide the vector into these partitions and do recursion
        partition1 = vector1[median1:]
        partition2 = vector2[:median2+1]
    else:
        # in this case, we partition the vectors similarly, but in the opposite
        # directions
        partition1 = vector1[:median1+1]
        partition2 = vector2[median2:]
    # recursion step
    return find_median(partition1, partition2, len(partition1))
```

Since we are only querying the middle element of the vector in each recursion step; and deciding which direction we need to partition the vector. We are effectively doing binary search in both vectors. Therefore the total number of queries are $O(\log n)$.